


STAT406- Methods of Statistical Learning Lecture 7

Matias Salibian-Barrera

UBC - Sep / Dec 2017



<http://xkcd.com/552/>

Model / feature selection - LASSO

- Another regularized method is given by LASSO

$$\min_{\alpha, \beta} \sum_{i=1}^n (y_i - \alpha - \beta' \mathbf{x}_i)^2 + \lambda \sum_{j=1}^p |\beta_j|$$

$$\min_{\alpha, \beta} \sum_{i=1}^n (y_i - \alpha - \beta' \mathbf{x}_i)^2 + \lambda \|\beta\|_1$$

for some $\lambda > 0$

Model / feature selection - LASSO

- The above is equivalent to

$$\min_{\alpha, \beta} \sum_{i=1}^n (y_i - \alpha - \beta' \mathbf{x}_i)^2$$

subject to

$$\sum_{j=1}^p |\beta_j| \leq K$$

for some $K > 0$

LASSO


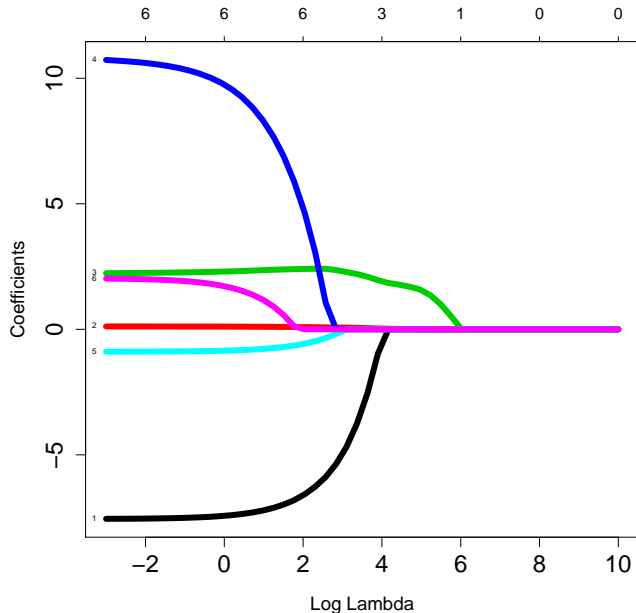


FIGURE 3.11. *Estimation picture for the lasso (left) and ridge regression (right). Shown are contours of the error and constraint functions. The solid blue areas are the constraint regions $|\beta_1| + |\beta_2| \leq t$ and $\beta_1^2 + \beta_2^2 \leq t^2$, respectively, while the red ellipses are the contours of the least squares error function.*

Credit data - glmnet output



Credit data - glmnet output

```
a <- glmnet(x=xm, y=yc, lambda=lambdas,  
            family='gaussian', alpha=1, intercept=FALSE)  
  
> coef(a, s=1)  
7 x 1 sparse Matrix of class "dgCMatrix"  
1  
(Intercept) .  
Income      -7.4285710  
Limit       0.1078894  
Rating      2.3006418  
Cards       9.7499618  
Age        -0.8515917  
Education   1.7182477
```

Credit data - glmnet output

```
> coef(a, s=exp(4))  
7 x 1 sparse Matrix of class "dgCMatrix"  
1  
(Intercept) .  
Income      -0.63094341  
Limit       0.02749778  
Rating      1.91772580  
Cards       .  
Age         .  
Education   .
```


Credit data - another implementation

```
> library(lars)
> b <- lars(x=xm, y=yc, type='lasso', intercept=FALSE)
> coef(b)
```

	Income	Limit	Rating	Cards	Age	Education
[1,]	0.0000000	0.000000000	0.000000	0.000000	0.0000000	0.000000
[2,]	0.0000000	0.000000000	1.835963	0.000000	0.0000000	0.000000
[3,]	0.0000000	0.01226464	2.018929	0.000000	0.0000000	0.000000
[4,]	-4.703898	0.05638653	2.433088	0.000000	0.0000000	0.000000
[5,]	-5.802948	0.06600083	2.545810	0.000000	-0.3234748	0.000000
[6,]	-6.772905	0.10049065	2.257218	6.369873	-0.6349138	0.000000
[7,]	-7.558037	0.12585115	2.063101	11.591558	-0.8923978	1.998283

```
> b
```

Call:


```
lars(x = xm, y = yc, type = "lasso", intercept = FALSE)
```

R-squared: 0.878

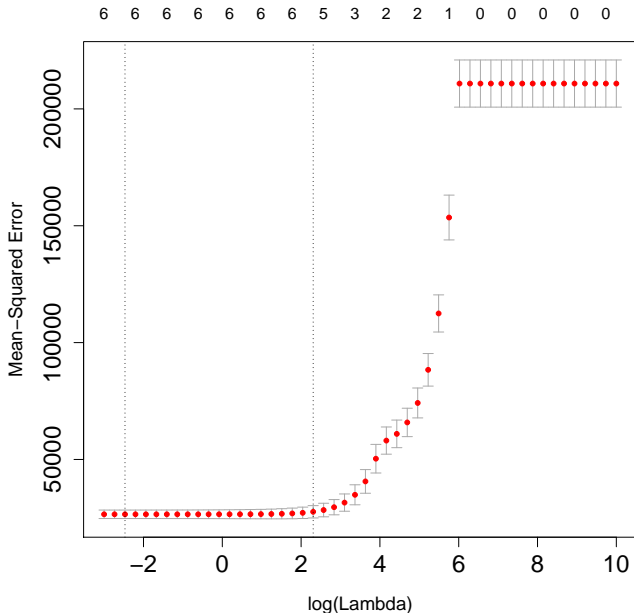
Sequence of LASSO moves:

	Rating	Limit	Income	Age	Cards	Education
Var	3	2	1	5	4	6
Step	1	2	3	4	5	6


Credit data - lars output



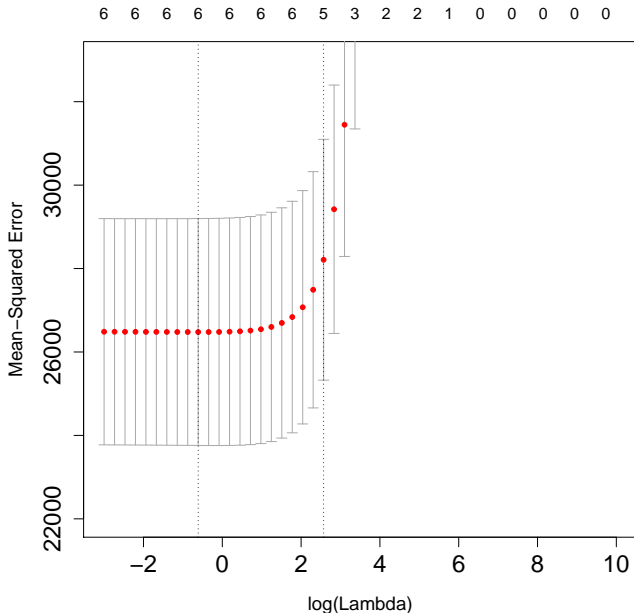
Credit data - CV - glmnet



Credit data - CV - another run




Credit data - CV - zoom



Model / feature selection - LASSO

Credit – 10 runs 5-fold CV



Model / feature selection - LASSO

- Worse estimated MSPE than Ridge Regression in this case
- It provides a sequence of explanatory variables, an ordered set of models
- Much like stepwise, but with better MSPE in this case

Model / feature selection - LASSO

- Why does it work? It is the convex proxy for the “nuclear norm”
- Also generates infinitely many estimates, but there’s a clever algorithm
- Inference?

Model / feature selection - LASSO

- When covariates are correlated, LASSO will typically pick any one of them, and ignore the rest
- Ridge Regression, on the other hand, combines the coefficients of correlated covariates, but doesn't provide sparse models

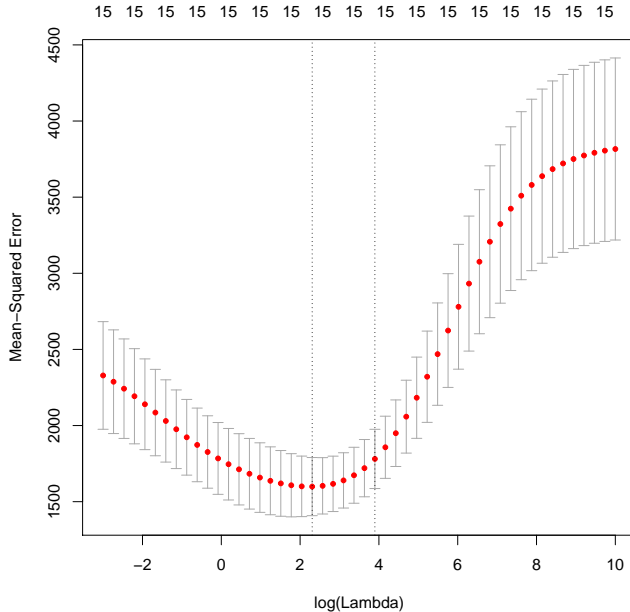
Ridge vs. LASSO

- Compare Ridge and LASSO on the air pollution data

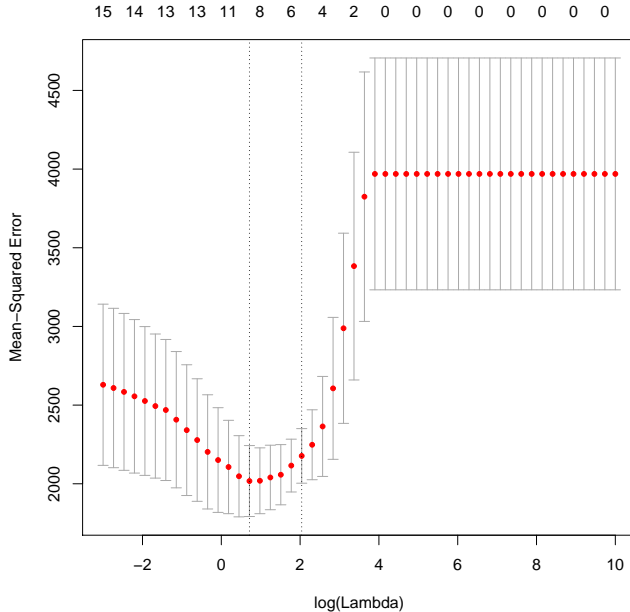
Air pollution example

```
airp <- read.table('../-30861_CSV-1.csv',  
  header=TRUE, sep=',')  
y <- as.vector(airp$MORT)  
xm <- as.matrix(airp[, names(airp) != 'MORT'])  
# Ridge  
set.seed(123)  
air.l2 <- cv.glmnet(x=xm, y=y, lambda=lambdas,  
  nfolds=5, alpha=0, family='gaussian',  
  intercept=TRUE)  
# LASSO  
set.seed(23)  
air.l1 <- cv.glmnet(x=xm, y=y, lambda=lambdas,  
  nfolds=5, alpha=1, family='gaussian',  
  intercept=TRUE)
```

Air pollution - Ridge



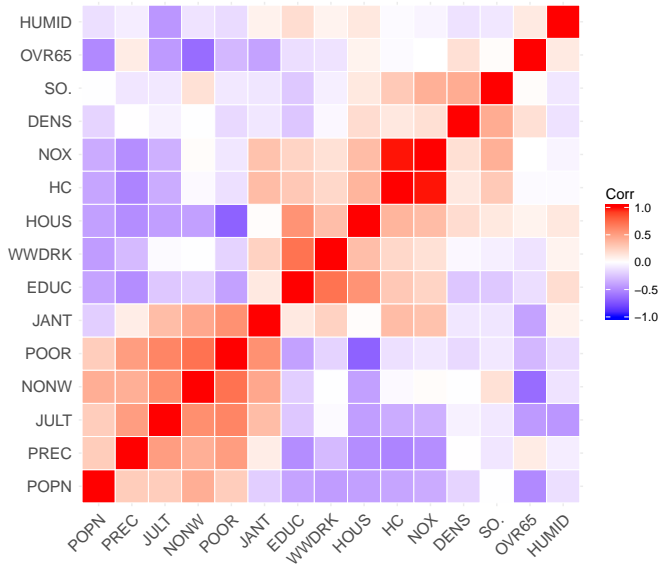
Air pollution - LASSO



Air pollution example

	Ridge	LASSO
(Intercept)	1179.335	1100.355
PREC	1.570	1.503
JANT	-1.109	-1.189
JULT	-1.276	-1.247
OVR65	-2.571	.
POPN	-10.135	.
EDUC	-8.479	-10.510
HOUS	-1.164	-0.503
DENS	0.005	0.004
NONW	3.126	3.979
WWDK	-0.476	-0.002
POOR	0.576	.
HC	-0.035	.
NOX	0.064	.
SO.	0.240	0.228
HUMID	0.372	.

Air pollution - Correlations



Model / feature selection - LASSO

- Oracle - consistency
- Problem: when $n < p$, LASSO will only choose up to n variables
- When covariates are correlated, LASSO will typically pick any one of them, and ignore the rest
- Ridge Regression, on the other hand, combines the coefficients of correlated covariates, but doesn't provide sparse models

Elastic Net

- Elastic Net is a compromise between the two:

$$\min_{\beta_0, \beta} \sum_{i=1}^n (y_i - \beta_0 - \beta' \mathbf{x}_i)^2 + \lambda \left[\alpha \|\beta\|_1 + \frac{(1 - \alpha)}{2} \|\beta\|_2^2 \right]$$

for some $\lambda > 0$ and $0 \leq \alpha \leq 1$.

Elastic Net

- $\alpha = 0$ reduces to Ridge Regression
- $\alpha = 1$ reduces to LASSO
- α needs to be chosen... how would you find a good choice for α ?

Air pollution example

- There are correlated covariates
- LASSO solution picks one of each group early on and relegates the rest to the end of the sequence
- Ridge Regression includes all variables always
- EN with $\alpha = 0.10$ gives a nice path of solutions...
- CV? bivariate search, unless α can be chosen beforehand

More flexible regression


- What if the regression function

$$E[Y|\mathbf{X}] = f(\mathbf{X})$$

is not linear?

- Example LIDAR

LIDAR



Non-linear regression

- Model: $E[Y|X_1, X_2, \dots, X_p] = f(X_1, X_2, \dots, X_p; \theta_1, \theta_2, \dots, \theta_k)$
- This is typically a non-linear model
- But it is fully parametric
- The parameters are $\theta_1, \theta_2, \dots, \theta_k$
- Using MLE (or LS) we can obtain estimates $\hat{\theta}_1, \dots, \hat{\theta}_k$
- ... and associated standard errors!

Non-linear regression

- Sometimes it's difficult to find an appropriate family of functions
- Polynomials are a natural choice

$$m(x) = m(x_0) + \frac{1}{2}m'(x_0)(x - x_0) + \dots$$
$$+ \frac{1}{k!}m^{(k-1)}(x_0)(x - x_0)^{k-1} + R_k$$


Non-linear regression

- Hence, we can try


$$E[Y|X] = \beta_0 + \beta_1 X + \beta_2 X^2 + \dots + \beta_k X^k$$

- This is a linear model! (**WHY?**)

LIDAR - 4th deg. polynomial



LIDAR - 10th deg. polynomial



More flexible bases

- Consider the (family) of function(s)

$$f_j(x) = (x - \kappa_j)_+ = \begin{cases} x - \kappa_j & \text{if } x - \kappa_j > 0 \\ 0 & \text{otherwise} \end{cases}$$

where κ_j are *knots*

- Model

$$E[Y|X] = \beta_0 + \beta_1 X + \sum_{j=1}^K \beta_{j+1} f_j(X)$$

- This is a linear model

More flexible bases


- The **knots** can be chosen arbitrarily
- It is customary to select them based on the sample

$$\kappa_j = \frac{j}{K+1} \text{ 100\% quantile of } x$$

- For example, with $K = 4$:

$$\kappa_1 = 20\%, \quad \kappa_2 = 40\%, \quad \text{etc.}$$

Regression splines, 5 knots



More flexible bases

- Consider a smoother basis


$$f_j(x) = (x - \kappa_j)_+^2 = \begin{cases} (x - \kappa_j)^2 & \text{if } x - \kappa_j > 0 \\ 0 & \text{otherwise} \end{cases}$$

where κ_j , $1 \leq j \leq K$ are *knots*


- Model

$$E[Y|X] = \beta_0 + \beta_1 X + \beta_2 X^2 + \sum_{j=1}^K \beta_{j+2} f_j(X)$$


Quadratic splines, 5 knots



Quadratic splines, 10 knots



Quadratic splines, 50 knots



More flexible bases

- Cubic splines will be useful


$$f_j(x) = (x - \kappa_j)_+^3 = \begin{cases} (x - \kappa_j)^3 & \text{if } x - \kappa_j > 0 \\ 0 & \text{otherwise} \end{cases}$$

where κ_j , $1 \leq j \leq K$ are *knots*


- Model

$$E[Y|X] = \beta_0 + \beta_1 X + \beta_2 X^2 + \beta_3 X^3 + \sum_{j=1}^K \beta_{j+3} f_j(X)$$


Cubic splines, 5 knots



Cubic splines, 10 knots




Cubic splines, 50 knots



More flexible bases

- Need to choose number and location of knots
- Need to make them less wiggly at the ends (Natural cubic splines)

Natural cubic spline, 5 knots




Smoothing splines

- Consider the following problem


$$\min_f \sum_{i=1}^n (Y_i - f(X_i))^2 + \lambda \int \left(f^{(2)}(t) \right)^2 dt$$

- The solution is a *natural* cubic spline with n knots at X_1, X_2, \dots, X_n .
- Natural* cubic splines are cubic splines with the restriction that they are linear beyond the boundary knots.


Smoothing spline, $\lambda = 0.20$




Smoothing spline, $\lambda = 0.50$



Smoothing spline, $\lambda = 0.75$



Smoothing spline, $\lambda = 2.00$



Selecting the penalty parameter

- How do we select λ ?
- Minimizing

$$RSS(\lambda) = \sum_{i=1}^n (Y_i - \mathbf{x}_i' \boldsymbol{\beta}_\lambda)^2$$

is not a good idea...

Selecting the penalty parameter

- Cross-validation: consider


$$CV(\lambda) = \sum_{i=1}^n \left(Y_i - \mathbf{x}_i' \boldsymbol{\beta}_{\lambda}^{(-i)} \right)^2$$

where $\boldsymbol{\beta}_{\lambda}^{(-i)}$ is the fit without using the point (Y_i, X_i)


and choose a value λ_0 such that

$$CV(\lambda_0) \leq CV(\lambda) \quad \forall \lambda \geq 0$$

5-fold CV, smoothing spline



Optimal fit via 5-fold CV



Selecting the penalty parameter

- Computing leave-one-out CV

$$CV(\lambda) = \sum_{i=1}^n \left(Y_i - \mathbf{x}'_i \boldsymbol{\beta}_{\lambda}^{(-i)} \right)^2$$

We might need to re-fit the model n times

Selecting the penalty parameter

- For some smoothers and models this is not necessary. For many linear smoothers $\hat{\mathbf{Y}} = \mathbf{S}_\lambda \mathbf{Y}$ we have

$$\hat{\mathbf{Y}}_r = \sum_{i=1}^n \mathbf{S}_{\lambda,r,i} Y_i \quad r = 1, \dots, n$$

and then

$$\hat{\mathbf{Y}}_r^{(-r)} = \frac{\sum_{i \neq r} \mathbf{S}_{\lambda,r,i} Y_i}{\sum_{i \neq r} \mathbf{S}_{\lambda,r,i}}$$

Selecting the penalty parameter

- Furthermore


$$\mathbf{S}_\lambda \mathbf{1} = \mathbf{1}$$

thus


$$\hat{\mathbf{Y}}_r^{(-r)} = \frac{\sum_{i \neq r} \mathbf{S}_{\lambda,r,i} Y_i}{1 - \mathbf{S}_{\lambda,r,r}}$$

$$CV(\lambda) = \sum_{i=1}^n \left(\frac{Y_i - \hat{\mathbf{Y}}_i}{1 - \mathbf{S}_{\lambda,i,i}} \right)^2$$

Optimal fit via leave-1-out CV



Compare with 5-fold CV optimal



Selecting the penalty parameter

- Computing $\mathbf{S}_{\lambda,i,i}$, $i = 1, \dots, n$ can be demanding

$$\begin{aligned} GCV(\lambda) &= \sum_{i=1}^n \left(\frac{Y_i - \hat{\mathbf{Y}}_i}{1 - \text{tr}(\mathbf{S}_{\lambda})/n} \right)^2 = \\ &= \frac{\sum_{i=1}^n (Y_i - \hat{\mathbf{Y}}_i)^2}{(1 - \text{tr}(\mathbf{S}_{\lambda})/n)^2} \end{aligned}$$