STAT406- Methods of Statistical Learning Lecture 18

Matias Salibian-Barrera

UBC - Sep / Dec 2017

Ensembles

- Ensembles of classifiers

- Combine classifiers trained on the same (or similar [e.g. bootstrapped]) data

- Consensus is reached by (equally weighted) voting or averaged estimated probabilities.

- Bagging and Random Forests are examples of ensembles.

Boosting

- Originally proposed for classification

- Main idea: sequentially re-train a simple classifier assigning more importance to points that were previously misclassified

Boosting

- The end result is a weighted average of all the classifiers

- Interesting ideas:

    - Not all components of the ensemble are treated equally

    - Members of the ensemble use information about other members

    - The underlying loss function has a "margin" (unlike 0-1 losses)

Boosting - AdaBoost.M1

Algorithm. Data $(y_i, \mathbf{x}_i)$, with $y_i \in \{-1, 1\}$

- Set initial weights $w_i = 1/n$, $1 \leq i \leq n$

- For $j = 1, \ldots, K$

- Build a classifier $T_j(\mathbf{x})$ to the data using weights w_i, $1 \leq i \leq n$

Boosting - AdaBoost.M1

- Let

$$e_{\mathbf{j}} = \left. \sum_{\mathbf{i}=1}^{n} w_{\mathbf{i}} I\left(y_{\mathbf{i}} \neq T_{\mathbf{j}}(\mathbf{x_i})\right) \middle/ \sum_{\ell=1}^{n} w_{\ell} \right.$$

- Let $\alpha_{\mathbf{j}} = \log\left((1 - e_{\mathbf{j}})/e_{\mathbf{j}}\right)$ and

$$w_{\mathbf{i}} = w_{\mathbf{i}} \exp\left(\alpha_{\mathbf{j}} I\left(y_{\mathbf{i}} \neq T_{\mathbf{j}}(\mathbf{x_i})\right)\right), \ \mathbf{i} = 1, \ldots, n$$

- Final classifier:

$$T(\mathbf{x}) = \mathrm{sign}\left(\sum_{\mathbf{j}=1}^{K} \alpha_{\mathbf{j}} \, T_{\mathbf{j}}(\mathbf{x})\right)$$

Error evolution - Spam data

AdaBoost error Vs number of trees

Error evolution - Isolet

AdaBoost error Vs number of trees

Boosting

- Boosting is fitting an **additive model**
- ... using a **forward search algorithm**
- ... and a **specific loss function**

Boosting

- Think of classifiers of the form

$$G(x) = \sum_{j=1}^{K} \beta_j \, f(\mathbf{x}, \gamma_j)$$

where $f(\mathbf{x}, \gamma_j)$ are simple base classifiers (e.g. trees)

Boosting

- Given a data set $(y_i, \mathbf{x}_i)$, $i = 1, \ldots, n$

$$\min_{G} \sum_{i=1}^{n} L\left(y_i, G(\mathbf{x}_i)\right) =$$

$$= \min_{\boldsymbol{\beta}, \boldsymbol{\gamma}} \sum_{i=1}^{n} L\left(y_i, \sum_{j=1}^{K} \beta_j \, f(\mathbf{x}_i, \gamma_j)\right)$$

where $\boldsymbol{\beta} = (\beta_1, \ldots, \beta_K)'$ and
$\boldsymbol{\gamma} = (\gamma_1, \ldots, \gamma_K)'$

Boosting

- Find approximate solutions sequentially

- Start with $f_0(\mathbf{x}) = 0$

- `for(j in 1:K)`

- Find

$$(\beta_j, \gamma_j) = \arg\min_{\beta, \gamma} \sum_{i=1}^{n} L\big(y_i, f_{j-1}(\mathbf{x}_i) + \beta\, f(\mathbf{x}_i, \gamma)\big)$$

- Let $f_j(\mathbf{x}) = f_{j-1}(\mathbf{x}) + \beta_j\, f(\mathbf{x}, \gamma_j)$

Counterexample

```
> set.seed(123)
> n <- 100
> x1 <- rnorm(n)
> x2 <- rnorm(n)
> y <- 2 - x1 + 6*x2 + rnorm(n, sd=.7)
> m1 <- lm(y~x1+x2)
> ( obj1 <- sum( resid(m1)^2 ) )
[1] 43.01311
>
> r1 <- y - mean(y)
> m2 <- lm(r1~x1-1)
> r2 <- resid(m2)
> m3 <- lm(r2~x2-1)
> ( obj2 <- sum( resid(m3)^2 ) )
[1] 109.3264
```

13